Career development and transition funding

This program aims to provide support – in the form of funding for graduate study, unpaid internships, independent study, career transition and exploration periods, and other activities relevant to building career capital – for individuals at any career stage who want to pursue careers that could help to reduce global catastrophic risks or otherwise improve the long-term future.

Apply for funding here.

We're especially interested in supporting individuals who want to pursue careers that are in some way related to mitigating potential risks posed by future advances in artificial intelligence or global catastrophic biological risks.

Applications are open until further notice and will be assessed on a rolling basis.

Generally speaking, we aim to review proposals within 6 weeks of receiving them, although this may not prove possible for all applications. Candidates who require more timely decisions can indicate this in their application forms, and we may be able to expedite the decision process in such cases.

Until recently, this program was known as the "early-career funding program", but we've decided to broaden its scope to explicitly include later-career individuals.

Scope

We're open to receiving applications from individuals who are already pursuing careers related to reducing global catastrophic risk (or otherwise improving the long-term future), looking to transition into such careers from other lines of work, or only just starting their ca

Contents +

pe

therefore a correspondingly wide range of proposals we would consider funding.

career transition and exploration periods, postdocs, obtaining professional certifications, online courses, and other types of one-off career-capital-building activities.

To name a few concrete examples of the kinds of applicants we're open to funding, in no particular order:

- A final-year undergraduate student who wants to pursue a master's or a PhD program in machine learning in order to contribute to technical research that helps mitigate risks from advanced artificial intelligence.
- An individual who wants to do an unpaid internship at a think tank focused on biosecurity, with the aim of pursuing a career dedicated to reducing global catastrophic biological risk.
- A former senior ML engineer at an AI company who wants to spend six months on independent study and career exploration in order to gain context on and investigate career options in AI risk mitigation.
- An individual who wants to attend law school or obtain an MPP, with the aim of working in government on policy issues relevant to improving the long-term future.
- A recent physics PhD who wants to spend six months going through a self-guided ML curriculum and working on projects in interpretability, in order to transition to contributing to technical research that helps mitigate risks from advanced AI systems.
- A software engineer who wants to spend the next three months doing independent study in order to gain relevant certifications for a career in information security, with the longer-term goal of working for an organization focused on reducing global catastrophic risk.
- An experienced management consultant who wants to spend three months exploring different ways to apply their skill set to reducing global catastrophic risk and applying to relevant jobs, with an eye to transitioning to a related career.
- A PhD graduate in an unrelated sub-area of computational biology who wants to spend four months getting up to speed on DNA synthesis screening in order to transition to working on this topic.
- A professor in machine learning, theoretical computer science, or another technical field who wants funding to take a one-year sabbatical to explore ways to contribute to technical AI safety or AI governance.

Funding criteria

We are particularly interested in funding people who have deeply engaged with questions about global catastrophic risk and/or the long-term future, and who have skills and abilities that could allow them to make substantial contributions in the relevant areas.

- Candidates should describe how the activity for which they are seeking funding will help them enter or transition into a career path that plausibly allows them to make these contributions. We appreciate that candidates' plans may be uncertain or even unlikely to work out, but we are looking for evidence that candidates have thought in a critical and reasonably detailed manner about those plans — not just about what career path(s) might open up for them, but also about how entering said career path(s) could allow them to reduce global catastrophic risk or otherwise positively impact the long-term future.
- We are looking to fund applications where our funding would make a difference — i.e. where the candidate is otherwise unable to find sufficient funding, or the funding they were able to secure imposes significant restrictions or requirements on them (for example, in the case of graduate study, restrictions on their research focus or teaching requirements). We may therefore turn down promising applicants who were able to secure equivalent support from other sources.
- If you receive a grant from us through this program, we will ask that you notify us about any other income or funding (e.g. paid work, fellowships, grant income, etc.) that you receive during the grant period. If your funding relates to a degree program, we will also ask you to notify us of any changes to your enrollment status. Depending on the nature of any additional income/enrollment status changes, we may alter the grant amount or timeframe in accordance with this policy.
- The program is open to applicants in any country.

Other information

- There is neither a maximum nor a minimum number of applications we intend to fund; rather, we intend to fund any applications that seem sufficiently promising to us to be above our general funding bar for this program.
- In some cases, we may ask outside advisors to help us review and evaluate applications. By submitting your application, you agree that we may share your application with our outside advisors for evaluation purposes.
- We encourage individuals with diverse backgrounds and experiences to apply, especially self-identified women and people of color.
- We plan to respond to all applications.

applied to that program should apply to this program instead. (We may decide to split out the Biosecurity Scholarship again as a separate program at a later point, but for practical purposes, current applicants can ignore this.)

- We may make changes to this program from time to time. Any such changes will be reflected on this page.

Past recipients

Alexandra Souly

Alexandra was awarded a CDTF grant to support her transition from software engineering to AI safety research.

At the time, Alexandra was a software engineer at Microsoft working on cloud automation. After completing the BlueDot AGI Strategy safety course, she recognized the urgency of AI safety work and decided to transition careers. The grant allowed her to quit her job and pivot without financial uncertainty, funding her Master's degree in Machine Learning at University College London — her institute of choice — and enabling her to spend a summer upskilling to get the most out of her degree.

At UCL, Alexandra collaborated with PhD students at the DARK Lab on research related to cooperative AI and multi-agent reinforcement learning, co-authoring a NeurIPS workshop paper and contributing to a now prominent benchmark in the field.

After completing her Master's, Alexandra interned at the Center for Human-Compatible AI (CHAI) at UC Berkeley, where she worked on value learning before pivoting to LLM safety. There, she co-authored *StrongREJECT*, a paper that established one of the first jailbreak benchmarks and datasets, which remains widely used and cited.

Alexandra then joined the newly established UK AI Security Institute's Safeguards team (now the Red Team), where she has spent the past two years. In this role, she has developed internal evaluation suites, contributed to frontier model pre-deployment evaluations, and co-authored research including *AgentHarm*, an agentic prompt injection benchmark, and a data poisoning paper with Anthropic and the Alan Turing Institute. Today, Alexandra work focuses on alignment red teaming: investigating whether frontier models exhibit signs of misalignment.

Bryce Cai

Bryce was awarded a CDTF grant to support three to six months of career transition into technical AI safety research.

Bryce studied chemistry and mathematics as an undergraduate, drawn to work with the potential for real-world impact. In 2021, he co-founded a Y Combinator-backed startup making computational biology tools more accessible, helping researchers without programming expertise run powerful tools like AlphaFold in the cloud. Following a period of health challenges in 2023, Bryce re-evaluated his career path. Informed by resources like 80,000 Hours, he concluded that AI safety was among the most pressing problems he could work on.

through introductory projects on subjects like interpretability, control, and model organisms research, but switched gears after being offered an internship at SecureBio. With CG's support, Bryce was able to set aside his original plan and specialize. During his internship, he developed ABLE (Agentic Bio Tool LLM Evaluation), an evaluation suite assessing how well AI agents can independently use computational biology tools to design proteins — work with significant biosecurity implications.

The internship led to a full-time offer. Bryce now works full-time at SecureBio, contributing to biosecurity-relevant evaluations and safeguards used by frontier model developers, labs, and national governments. Reflecting on the grant, Bryce shared that it "took a large weight off [his] back". Without it, he estimates that his transition into AI safety would have been delayed by several months.

Robi Rahman

Robi was awarded a CDTF grant to support his transition from chemical engineering to AI safety research. After three years of earning to give, Robi recognized the potential impact of AI on humanity's future and resolved to redirect his career.

While completing his master's degree in data science from Harvard University, Robi received a CDTF grant to undertake four months of part-time AI self-study and internships. During this period, he gained hands-on experience in machine learning by developing an endangered wildlife identification algorithm for WildTrack, a conservation nonprofit, followed by a part-time role at the Stanford Institute for Human-Centered Artificial Intelligence, doing research and writing for the 2023 AI Index Report.

The grant allowed Robi to focus on technical skill development without financial strain, bridging the gap between his existing skills and those needed to work in AI safety.

Robi then applied those skills while working at Epoch AI as a data scientist from 2023-2025, tracking the algorithms, investments, and hardware used to develop advanced AI systems. He now works on the Technical Governance Team at MIRI.

Cecil Abungu

Cecil was awarded two CDTF grants that kickstarted his career in AI safety–related governance.

The first grant, awarded in early 2020, helped Cecil transition from his focus on constitutional law to AI governance research. This grant gave Cecil the chance to fully immerse himself in AI governance over the course of a year, with two key research periods: six months at the Institute for Law and AI (formerly the Legal Priorities Project) and another six months with the Centre for the Study of Existential Risk's AI:FAR team. Cecil delved deep into the field during this period, reading extensively, watching relevant lectures, and participating in research seminars as both an attendee and presenter. This exploration laid a strong foundation for his future work and connected him with leading experts.

AI governance research and policy engagement in Africa. “Without the grant, I would probably have to do other kinds of work to finance my PhD”, Cecil shared.

Looking forward, Cecil hopes that his research will help policymakers take AI risks more seriously. “I want to positively influence the thinking of other AI safety researchers and help more talented Africans begin researching AI safety–related governance in a consequential way.”

Application process

Applications are open until further notice. You can apply using this form.

Required application materials:

- Proposal, no longer than 250 words
- Personal statement, no longer than 250 words
- Approximate budget, no longer than half a page
- CV or resume, no longer than 2 pages
- Academic transcript (undergrad and graduate, if applicable)
- Answers to a few other questions (see application form)

We may contact you to request additional information. And in some cases, the assessment may also involve a brief interview, via video teleconference.

We are aware that if you are applying for graduate school or internships, you will typically not know at this point which specific programs will admit you. If you are applying to several different but related programs and are doing so with a similar career trajectory in mind (e.g. different law schools and MPP programs to pursue a career in public policy), please specify this in your proposal and budget and submit a single application. If you are applying to several rather different programs with clearly distinct career trajectories in mind (e.g. policy schools and history PhD programs), please submit separate applications.

If you have questions, you can reach us via careerdevelopmentfunding@coefficientgiving.org.

Coefficient

Careers

Media Center

Partner With Us

Contact Us

Privacy Policy

Governance

Sign up for news about our grants, research, and more

Email

Subscribe

X (Twitter)

LinkedIn